

Notebook: Geometric multigrid for 2D-PDE's

Solving Partial Differential Equations has never been an easy task on mathematicians. In fact, in many cases it's even considered *impossible* to do with our current mathematics, so our best way to approach it is with numerical/computational approximations. It's this exact field that has specialized in solving PDE's numerically: *Using Iterative methods, Fourier Transforms, Finite Differences, etc..*

This project is about the **Multigrid methods** which are one of the most interesting assets of the numerical computing toolbox. We're also going to do some benchmarking of their performance regarding their iterative counterparts.

Finally, to be able to start, we should first do the initial step of the numerical problem resolution framework: **Discretizing the PDE space**

```
• # Used Packages Index
• begin
•   using LinearAlgebra      # Linear Algebra
•   using SparseArrays      # Sparse Arrays Optimization
•   using Plots              # Plotting & Visualization
• end ;
```

Constructing the Laplacian 2D-operator and the Problems PDEs

$$Au = \Delta u = \frac{\partial^2 u}{\partial x^2} e_x + \frac{\partial^2 u}{\partial y^2} e_y$$

To construct the laplacian operator, we're using the **Kronecker product** with the identity matrix to develop the axe-wise differentiation operators and sum them up lately to obtain Δ .

$$\begin{bmatrix} A \otimes B_{11} & \dots & A \otimes B_{1m} \\ \vdots & \ddots & \vdots \\ A \otimes B_{m1} & \dots & A \otimes B_{mm} \end{bmatrix}$$

$$\begin{bmatrix} -4 & 1 & . & 1 & . & . & . & . & . \\ 1 & -4 & 1 & . & 1 & . & . & . & . \\ . & 1 & -4 & . & . & 1 & . & . & . \\ 1 & . & . & -4 & 1 & . & 1 & . & . \\ . & 1 & . & 1 & -4 & 1 & . & 1 & . \\ . & . & 1 & . & 1 & -4 & . & . & 1 \\ . & . & . & 1 & . & . & -4 & 1 & . \\ . & . & . & . & 1 & . & 1 & -4 & 1 \\ . & . & . & . & . & 1 & . & 1 & -4 \end{bmatrix}$$

Δ Operator Matrix

Regarding the Kronecker product, it is actually a **computationally expensive** operation, that is why we have to use **sparse matrices** for this extent. The use of sparse matrices is further supported by the sparsity factor of our Laplacian operator that isn't to be ignored.

We then need to transform our differential equation into a linear problem in the form of $\mathbf{A}\mathbf{u}=\mathbf{f}$. For that, we only have to use the matrix σI to then factorize the equation by $u(x, y)$.

$$\begin{aligned} -\Delta u(x, y) + \sigma u(x, y) &= f(x, y) \quad (P_1) \\ (\sigma I - \Delta)u(x, y) &= f(x, y) \\ A_1 &= \sigma I - \Delta \end{aligned}$$

As for the Anisotropic problem P_2 , the construction is similar but only using the partial derivatives

$$\begin{aligned} -\frac{\partial^2 u(x, y)}{\partial x^2} - \epsilon \frac{\partial^2 u(x, y)}{\partial y^2} &= f(x, y) \quad (P_2) \\ A_2 &= \left(\frac{\partial^2}{\partial x^2} - \epsilon \frac{\partial^2}{\partial y^2} \right) I \end{aligned}$$

```

• begin
•   function A1(n::Int, σ::Float64)
•       ∂² = Tridiagonal(ones(n-1), -2 * ones(n), ones(n-1))
•       ∂x² = (n²) * kron(sparse(∂²), I(n))      # kron is the Kronecker product
•       ∂y² = (n²) * kron(I(n), sparse(∂²))      # h is the unit of displacement
•       Δ = ∂x² + ∂y²
•       return σ * I(n²) - Δ
•   end
•
•   function A2(n::Int, ε::Float64)
•       ∂² = Tridiagonal(ones(n-1), -2 * ones(n), ones(n-1))
•       ∂x² = (n²) * kron(sparse(∂²), I(n))      # kron is the Kronecker product
•       ∂y² = (n²) * kron(I(n), sparse(∂²))      # h is the unit of displacement
•       return - ∂x² - ε * ∂y²
•   end
•
•   A1(σ::Float64) = n -> A1(n, σ)
•   A2(ε::Float64) = n -> A2(n, ε)
• end ;

```

Before jumping to start modelling the solvers, we need first to iterate on the boundaries conditions. The PDE's we're solving are both using the same boundaries conditions $u(x, y) = 0$ on $\partial\Omega$ so we're just coding them in a function that will be called after each solver iteration to reinject the boundaries so that we don't miss the track of them.

```

• function boundaries(v)
•     n = Int(sqrt(size(v,1)))
•     grid_v = reshape(v, (n, n))'
•     grid_v[1,:] .= 0
•     grid_v[end,:] .= 0
•     grid_v[:,1] .= 0
•     grid_v[:,end] .= 0
•     return grid_v |> transpose |> vec
• end ;

```

Constructing the solvers algorithms

For the stationary smoothers, we're each time choosing one of the defined functions below either to be used on its own for comparison or to use it as the smoother in our Multigrid implementation.

```

• function Jacobi(A, b, u₀ = zeros(size(A, 1)), ε = 1e-7, maxiter = 10, bounds =
false)
•     u = u₀; n = Int(sqrt(size(A, 1))); iter = 0
•     M = Diagonal(A)
•     N = UnitLowerTriangular(A) + UnitUpperTriangular(A) - 2*I(n^2)
•     while iter <= maxiter
•         iter += 1
•         if bounds
•             u = boundaries(u)
•         end
•         u = inv(M) * (N*u + b)
•         (norm(b - A*u, 2) > ε) || break
•     end
•     return u
• end ;

```

```

• function JOR(A, b, ω, u₀ = zeros(size(A, 1)), ε = 1e-7, maxiter = 10, bounds =
false)
•     u = u₀; iter = 0
•     M = Diagonal(A) / ω
•     while iter <= maxiter
•         iter += 1
•         if bounds
•             u = boundaries(u)
•         end
•         r = b - A*u
•         z = inv(M) * r
•         u += z
•         (norm(r, 2) > ε) || break
•     end
•     return u
• end ;


```

```

• function SOR(A, b, ω, u₀ = zeros(size(A, 1)), ε = 1e-7, maxiter = 10, bounds =
false)
•     u = u₀; n = Int(sqrt(size(A, 1))); iter = 0
•     D = Diagonal(A)
•     L = UnitLowerTriangular(A) - I(n^2)
•     U = UnitUpperTriangular(A) - I(n^2)
•     while iter <= maxiter
•         iter += 1
•         if bounds
•             u = boundaries(u)
•         end
•         u = inv(D + ω * L) * (ω * b - (ω * U + (ω-1) * D) * u)
•         (norm(b - A*u, 2) > ε) || break
•     end
•     return u
• end ;

```

Grid Spacing



As shown in the simple form of grid-spacing in a **V-cycle** multigrid, we need some **interpolation** operators to be used to move from one grid to another:

- *Restriction*: From a *fine* grid to a *coarser* one
- *Prolongation*: From a *coarse* grid to a *finer* one

Therefore, the methods defined below will be used interchangeably.

Also, as we're treating 2 different problems:

- *isotropic*: $A_1 u(x, y) = f(x, y)$
- ϵ -*anisotropic*: $A_2 u(x, y) = f(x, y)$

That's why we're going to need two different strategies for grids interpolation as we're not moving in the same space for both problems.

```

• # Multigrid's **Isotropic** Interpolation operators
• begin
•   # Restriction
•   injection(grid) = grid[2:2:end, 2:2:end]
•
•   function halfweight(grid)
•     g = Float64.(grid)
•     for i=2:2:size(grid,1)-1, j=2:2:size(grid,2)-1
•       g[i,j] = g[i,j] / 2 - (
•         g[i-1,j] + g[i+1,j]
•         + g[i,j-1] + g[i,j+1]) / 8
•     end
•     return injection(g)
•   end
•
•   function fullweight(grid)
•     g = Float64.(grid)
•     for i=2:2:size(grid,1)-1, j=2:2:size(grid,2)-1
•       g[i,j] = g[i,j] / 4 - (
•         g[i-1,j] + g[i+1,j]
•         + g[i,j-1] + g[i,j+1]) / 8 - (
•         g[i-1,j-1] + g[i+1,j+1]
•         + g[i-1,j+1] + g[i+1,j-1]) / 16
•     end
•     return injection(g)
•   end
• end

```

```

.
.
. # Prolongation
. function enlarge(grid)
.     n = size(grid,1) * 2
.     n == 2 && return repeat(grid, n, n)
.     g = zeros((n,n))
.     for i=2:n-1, j=2:n-1
.         g[i, j] = grid[i+2, j+2]
.     end
.     return g
. end
.
. function linearize(grid)
.     n = size(grid,1) * 2
.     n == 2 && return repeat(grid, n, n)
.     g = zeros((n,n))
.     for i=2:n-1, j=2:n-1
.         g[i, j] = (grid[Int(floor((i+1)/2)), Int(floor((j+1)/2))]
.             + grid[Int(ceil((i+1)/2)), Int(floor((j+1)/2))]
.             + grid[Int(floor((i+1)/2)), Int(ceil((j+1)/2))]
.             + grid[Int(ceil((i+1)/2)), Int(ceil((j+1)/2))]) / 4
.     end
.     return g
. end
. end ;

```

Algorithm 1: V-cycle multigrid iteration

$w_h := \text{MG}(h, L_h, w_h^0, b_h)$

```

IF ( $h == h_0$ )
    Solve exactly:  $w_h = L_h^{-1} b_h$ 
ELSE
    Pre-smooth  $v_1$  times:  $w_h := \text{SMOOTH}^{v_1}(L_h, w_h^0, b_h)$ 
    Restriction:  $r_H := I_h^H(b_h - L_h w_h)$ 
    Initialize correction:  $\delta_H := 0$ 
    Recursion :  $\delta_H := \text{MG}(H, L_H, \delta_H, r_H)$ 
    Prolongation:  $\delta_h := I_H^h \delta_H$ 
    Correction:  $w_h := w_h + \delta_h$ 
    Post-smooth  $v_2$  times:  $w_h := \text{SMOOTH}^{v_2}(L_h, w_h, b_h)$ 
ENDIF
RETURN  $w_h$ .

```

```

. """
.     multigrid(A, b, u, l, ω, ε, steps,
.             restrict, prolong, iter)
.
. l-level Multigrid cycle. Each sub-level dive is held with x10 smoothing iterations.
. - A: the PDE's linear operator constructor
. - restrict: function for grid restriction
. - prolong: function for grid prolongation
. - iter: number of iterations of current grid smoothing
. - steps: number of multigrid callbacks per cycle:
.     - 1 (default) is for V-cycle
.     - 2 is for normal W-cycle
.     - n (> 2) is for W-cycles with n sub-refinements
. """
. function multigrid(A, b, u, l, ω, ε=1e-7, steps=1,
.     restrict=injection, prolong=enlarge, iter=10)
.     n = Int(sqrt(size(b, 1)))
.     A_n = A(n)
.     if l == 0
.         # We can also use a direct solver instead

```

```

•     # u = Array(An) | b
•     u = JOR(An, b, ω, u, ε, iter)           # Resolution
• else
•     u = JOR(An, b, ω, u, ε, iter)           # Pre-smoothing
•
•     # Defect restriction
•     r = reshape(b - An*u, (n, n)) |> transpose |>
•         restrict |> transpose |> vec
•
•     # Coarse-level Correction
•     δr = zeros(size(r))
•     for i=1:steps
•         δr = multigrid(A, r, δr, l-1, ω, ε,
•             steps, restrict, prolong, iter*3)
•     end
•
•     # Defect Prolongation δr → δ
•     δ = reshape(δr, (n+2, n+2)) |> transpose |>
•         prolong |> transpose |> vec
•
•     u += δ                                   # Correction
•     u = JOR(An, b, ω, u, ε, iter)           # Post-smoothing
• end
• return u
• end ;

```

Simulation and Results Interpretation

We start by running our *isotropic* problem resolution governed by the A_1 operator. We're also only interested in the *unit square* resolution of the function f , and for it we discretize the space on $n = 1024$ units, $h = 1/n$ being the displacement unit.

```

• begin
•     n = 1024
•     h = 1 / n
•     ω = 0.95
•     σ = 0.7
•     A = A1(σ)
•     u = zeros(n^2)
•     f = 3
•     b = [sin(2π*f*i*j) for i=0:h:1-h for j=0:h:1-h]
• end ;

```

9.14521813

```

• # Jacobi Over Relaxation
• @elapsed u1 = JOR(A(n), b, ω, u, 1e-30, 50, true)

```

13.729009692

```

• # Multigrid V-cycle
• @elapsed u2 = multigrid(A, b, u, 3, ω, 1e-30, 1, injection, linearize)

```

Examining the cells above, we can notice that both solvers took a similar time to finish their iterations. If we then thoroughly examine the written multigrid code, we'll find out that for the current simulation, while it yielded nearly the **same execution-time** as the Jacobi-Over-Relaxation method, we're actually running a *3-level V-cycle multigrid* that computes over **800 iterations** and that also have given a by-far **better approximation** than its iterative counterpart even when running most of those iterations on the *coarser* grid 128x128 instead of the original 1024x1024 one.

"||e_m|| < ||e_j||"


"Error norm gain: 318.32821"

The near-similarity of the execution-time between both solvers is actually due to the fact that the 300 iterations ran by the Multigrid are equivalent to only 48 ones of the Jacobi-Over-Relaxation method.


"Multigrid JOR-relative cost: 48.359375 iterations"

Solutions Visualization & interpretation

JOR Solution Error




Multigrid Solution Error




From the graphs shown above and below, we can tell that the multigrid's solution is by far a better solution than the iterative JOR solver (*Multigrid's being dimmer and softer which by the graph's legend means a **nearer to 0 error***). Also, the solutions difference distribution tells us even more about the struggle that Jacobi-Over-Relaxation method had trying to polish the *lowest frequencies* in only 50 iterations.


Solutions Difference Heatmap



Solutions Difference Distribution



V-cycle vs W-cycle



We'll now proceed to compare the convergence of the V-cycle and the W-cycle variants of the multigrid solver for various σ and n levels for the problem operator $A_1(\sigma, n)$.

$n = 1024$, $\sigma = 0.7$

```

• begin
•    $\sigma_1 = 0.7$ 
•    $n_1 = 1024$ 
•    $h_1 = 1 / n_1$ 
•    $u_{n1} = \text{zeros}(n_1^2)$ 
•    $b_{n1} = [\sin(2\pi * i * j) \text{ for } i=0:h_1:1-h_1 \text{ for } j=0:h_1:1-h_1]$ 
•    $A_{11} = A_1(\sigma_1)$ 
• end ;

```

11.645717189

```

• # Multigrid V-cycle
• @elapsed  $u_{i1} = \text{multigrid}(A_{11}, b_{n1}, u_{n1}, 3, \omega, 1e-30, 1)$ 

```

26.908595232

```

• # Multigrid 2-steps W-cycle
• @elapsed  $u_{j1} = \text{multigrid}(A_{11}, b_{n1}, u_{n1}, 3, \omega, 1e-30, 2)$ 



```

By just executing both solvers, we can already notice the difference in time between both solvers. The difference shown above actually makes sense: Our W-cycle is a **2-step** one on a **3-level** multigrid, which means it takes the exact form of a W, then we're executing the same number of iterations but **twice**. This gets reflected by the *more-than-twice* execution time we have with the W-cycle, compared to the V-cycle.

" $\|e_v\| > \|e_x\|$ "

From here, we can confirm that the W-cycle was actually worth the wait, as it did help smoothing out the error in a more extreme way than the V-cycle did.

"Error norm gain: 95.90758"

V-cycle - $n = 1024$ - $\sigma = 0.7$ W-cycle - $n = 1024$ - $\sigma = 0.7$ 

2- $n = 1024$, $\sigma = 7$

```

• begin
•    $\sigma_2 = 7.$ 
•    $A_{12} = A_1(\sigma_2)$ 
•    $u_{i2} = \text{multigrid}(A_{12}, b_{n1}, u_{n1}, 3, \omega, 1e-30, 1)$       # Multigrid V-cycle
•    $u_{j2} = \text{multigrid}(A_{12}, b_{n1}, u_{n1}, 3, \omega, 1e-30, 2)$       # Multigrid W-cycle
•    $e_{i2} = b_{n1} - A_{12}(n) * u_{i2}$                                 # V-cycle solution error
•    $e_{j2} = b_{n1} - A_{12}(n) * u_{j2}$                                 # W-cycle solution error
•    $\text{fastr}_2 = \text{norm}(e_{i2}, 2) - \text{norm}(e_{j2}, 2)$ 
• end ;

```

" $||e_v|| > ||e_x||$ "

"Error norm gain: 103.13397"

"Error norm gain - V-cycle: 8.03361"

"Error norm gain - W-cycle: 15.26"

Increasing σ by a factor of x10 lead both solvers to better solutions relatively, but it did have the same effect on both of them as the error gain of the W-cycle over the V-cycle is still nearly the same.

3- $n = 1024$, $\sigma = 70$

```

• begin
•    $\sigma_3 = 70.$ 
•    $A_{13} = A_1(\sigma_3)$ 
•    $u_{i3} = \text{multigrid}(A_{13}, b_{n1}, u_{n1}, 3, \omega, 1e-30, 1)$       # Multigrid V-cycle
•    $u_{j3} = \text{multigrid}(A_{13}, b_{n1}, u_{n1}, 3, \omega, 1e-30, 2)$       # Multigrid W-cycle
•    $e_{i3} = b_{n1} - A_{13}(n) * u_{i3}$                                 # V-cycle solution error
•    $e_{j3} = b_{n1} - A_{13}(n) * u_{j3}$                                 # W-cycle solution error
•    $\text{fastr}_3 = \text{norm}(e_{i3}, 2) - \text{norm}(e_{j3}, 2)$ 
• end ;

```

" $||e_v|| > ||e_x||$ "

"Error norm gain: 121.73866"

"Error norm gain - V-cycle: 67.6177"

"Error norm gain - W-cycle: 86.22238"

Increasing σ by a factor of x100 lead both solvers to better solutions, and this time, it started showing its effect on the W side of things more than on the V-cycle's side. But still, it also increased the error-wise gain of the V-cycle multigrid solver too.

4- $n = 256$, $\sigma = 7$

```
• begin
•   n2 = 256
•   h2 = 1 / n2
•   un2 = zeros(n2^2)
•   bn2 = [sin(2*pi*f*i*j) for i=0:h2:1-h2 for j=0:h2:1-h2]
•   ui4 = multigrid(A12, bn2, un2, 3, w, 1e-30, 1) # Multigrid V-cycle
•   uj4 = multigrid(A12, bn2, un2, 3, w, 1e-30, 2) # Multigrid W-cycle
•   ei4 = bn2 - A12(n2)*ui4 # V-cycle solution error
•   ej4 = bn2 - A12(n2)*uj4 # W-cycle solution error
•   fastr4 = norm(ei4, 2) - norm(ej4, 2)
• end ;
```

"||e_v|| > ||e_x||"

"Error norm gain: 5.50434"

"Error norm gain - V-cycle: 417.54523"

"Error norm gain - W-cycle: 319.9156"

Decreasing n by a factor of x4 lead both solvers as usual to **better solutions**, this time, having a **huge effect on the V-cycle's** side compared to the W-cycle as the difference between the solutions is converging to none. Still, the latter is the better solver on these simulation settings.

As for the effect this had on improving the approximation, this was expected: Actually, as we go down in the grid's size, we're releasing the multigrid from the burden of having to interpolate on bigger grids (*the ones that we got rid of when we divided the fine size by 4*) which means the solver now is **less error-prone** than on the bigger grid case. Also, let's not forget that by doing such, we're giving the multigrid the opportunity to **spend the same number of iterations but on coarser grids**, this helps it converge faster than before. And this is exactly why we witnessed these huge evolutions on the approximation error.

5- $n = 32$, $\sigma = 7$

```
• begin
•   n3 = 32
•   h3 = 1 / n3
•   un3 = zeros(n3^2)
•   bn3 = [sin(2*pi*f*i*j) for i=0:h3:1-h3 for j=0:h3:1-h3]
•   ui5 = multigrid(A12, bn3, un3, 3, w, 1e-30, 1) # Multigrid V-cycle
•   uj5 = multigrid(A12, bn3, un3, 3, w, 1e-30, 2) # Multigrid W-cycle
```

```

•      ei5      = bn3 - A12(n3)*ui5           # V-cycle solution error
•      ej5      = bn3 - A12(n3)*uj5           # W-cycle solution error
•      fastr5   = norm(ei5, 2) - norm(ej5, 2)
• end ;

```

```
"||ev|| > ||ex||"
```

```
"Error norm gain: 0.00611"
```

```
"Error norm gain - V-cycle: 480.95131"
```

```
"Error norm gain - W-cycle: 377.82345"
```

On the abstract side of the story, as the theory states, the W-cycle is still a better solver than the V-cycle. But as we're still approaching coarser grid sizes (*here dividing by a factor of x32*), the fine-tuning is happening on even coarser grids (*with a size of 4x4 in these settings*) so increasing the number of iterations or further smoothing out the approximation n-times (*as it's the case for the W-cycle*) **isn't going to make any difference**.

Restriction Operators

Trying to visualize the workings of the Multigrid solver, we always stumble on the matrix transformations we're using to interpolate the grid we have to a coarser one (*restriction*) or to a finer one (*prolongation*). As we already created these 3 operators (*see code in the solvers implementation part*), we're going to compare them according to what they give in term of approximation error.

For this benchmarking, we're conducting the $n = 1024$, $\sigma = 7$ simulation. We're also going to use the **4-level 2-steps W-cycle** multigrid equipped with the **linearization** prolongation operator for better results and more complexity to be able to well-assess the performance difference between all runs.

```
45.967655775
```

```

• # Restriction by Injection
• @elapsed uj6 = multigrid(A12, bn1, un1, 4, ω, 1e-30, 2, injection, linearize)

```

```
46.039371688
```

```

• # Restriction by Halfweighting
• @elapsed uj7 = multigrid(A12, bn1, un1, 4, ω, 1e-30, 2, halfweight, linearize)

```

```
46.10548316
```

```

• # Restriction by Fullweighting
• @elapsed uj8 = multigrid(A12, bn1, un1, 4, ω, 1e-30, 2, fullweight, linearize)

```

From the execution-times of the 3 variants, we can conclude that thanks to Julia's loops optimization, we can use the 3 restriction operators interchangeably without any loss in performance.

```

• # Comparing solutions error rates
• begin
•      ej6      = bn1 - A12(n)*uj6           # Injection solution error - ei
•      ej7      = bn1 - A12(n)*uj7           # Halfweighting solution error - eh

```

```


•   ej8      = bn1 - A12(n)*uj8           # Fullweighting solution error - eφ
•   fastr67 = norm(ej6, 2) - norm(ej7, 2)
•   fastr68 = norm(ej6, 2) - norm(ej8, 2)
•   fastr78 = norm(ej7, 2) - norm(ej8, 2)
•   end ;

```


"||e_h|| < ||e_φ||"

"Error norm gain - HW/FW: 113.10642"

Approx. Difference - Halfweighting/Fullweighting




Approx. Difference Distribution




"||e_i|| < ||e_h||"

"Error norm gain - I/HW: 370.85025"

Approx. Difference - Injection/Halfweighting



Approx. Difference Distribution



"||e_i|| < ||e_φ||"

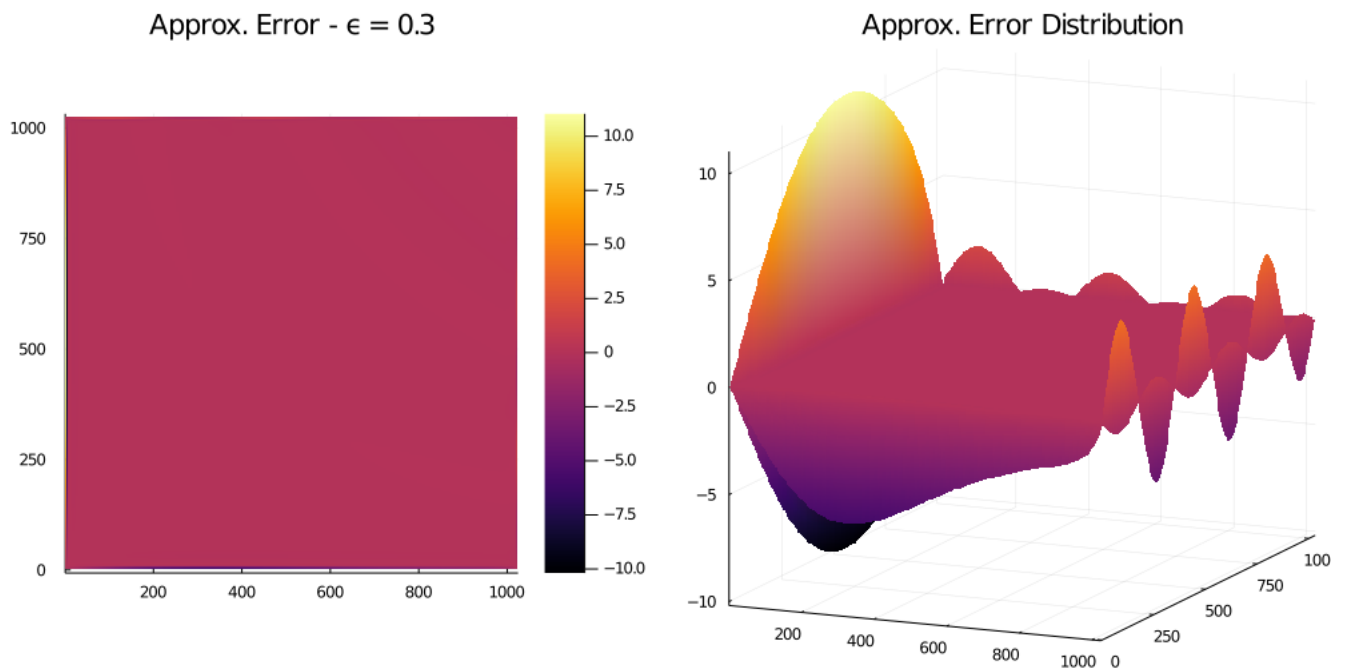
"Error norm gain I/FW: 483.95667"

What we learn from these results is that as long as we're not changing the local values of the solution matrices, there's no problem coarsening the grid some more. Which is in a bit the opposite result of the comparison between the *linearization* and the direct *enlargement* prolongations operators. This may be explained by the fact that when using the multigrid solver, we're trying to port our problem into coarser grids to solve it on that level, so if we're changing the values **even if it's a local-aware change** (as it's the case for *halfweighting* and even more in *fullweighting*), we may be **changing the target distribution** and by such furthering our path to the optimal minima.

Anisotropic Problem Resolution


We already constructed the Anisotropic problem PDE (P_2) in the PDE's & Operators Construction part of the document, so we're going to directly try to solve it using our 4-level 2-step W-cycle multigrid equipped by the injection & linearization interpolation operators.

1- $\epsilon = 0.3$




2- $\epsilon = 3$

Approx. Error - $\epsilon = 3$




Approx. Error Distribution




3- $\epsilon = 30$

Approx. Error - $\epsilon = 30$



Approx. Error Distribution



Trying to analyze the figures we get from the 3 simulations we ran, we can identify one main problem which is the **non-convergence on the pure directions** being X and Y axis of the vector space. Also, when we look into the difference between the first and the second graphs, we may conclude that the error is also *alternating from one direction into the other* while moving through intervals of values of ϵ .

This is actually caused by the nature of the **deformation** that the PDE applies to the vector space. As the equation P_2 states, differentiating the solution on the X-axis by one unit will directly force differentiating it by ϵ units onto the Y-axis.

$$-\frac{\partial^2 u(x, y)}{\partial x^2} - \epsilon \frac{\partial^2 u(x, y)}{\partial y^2} = f(x, y) \quad (P_2)$$

While this may sound not-harmful at all for our solvers, what we can guess is causing this issue, is actually our coarsening strategies. They're **based on equally distributed spaces** and that's why they're affecting **each one of the neighbors equal coefficients** when using neighbors-based coarsening. And this is the exact reason why the injection is yielding the best restriction therefore the best approximation among all the other operators on this problem.


One new way we can think of solving this problem, is to try to create a *halfweighting-like version of a restriction operator*, but **taking into account the ϵ parameter for the Y direction**.

```


• begin
•     """
•         e_halfweight(epsilon, grid)
•
•     A function that returns the epsilon-halfweighting of the grid taken as input.
•     """
•     function e_halfweight(epsilon, grid)
•         g = Float64.(grid)
•         for i=2:2:size(grid,1)-1, j=2:2:size(grid,2)-1
•             g[i,j] = epsilon * g[i,j] / 4 - (
•                 g[i-1,j] + g[i+1,j]
•                 + (1 - epsilon/2) * g[i,j-1] + (1 - epsilon/2) * g[i,j+1]) / 8
•         end
•         return injection(g)
•     end
•
•     """
•         e_halfweight(epsilon)
•
•     A function that constructs the epsilon-halfweighting operator.
•     """
•     e_halfweight(epsilon) = grid -> e_halfweight(epsilon, grid)
• end ;

```

ϵ -halfweighting - $\epsilon = 3$



Approx. Error Distribution



- $\text{norm}(\mathbf{e}_{j_{12}}, 2) - \text{norm}(\mathbf{e}_{j_{10}}, 2)$

And voilà!

We finally obtained an operator that worked better and produced a cleaner approximation with less error norm than the best operator among the restrictors tried on the anisotropic problem.

The idea behind it is that as the Y-axis has got **stretched out by** ϵ , then at every point of our grid, we'll have an ϵ -contribution of one element, and $1 - \epsilon/N$, N being the number of neighbors around it that are taken into account.

By such, we could decrease our norm by **55-units** on the euclidian norm of the vector space.